
XGBoost: Code Presentation

— Olivier Rochoaix, Erick Sanmartin, —
Logan Sell, Nolan Thomas

Data Processing - songs.csv

- Dropped unrelated columns
- Dropped rows for unique genres that skewed data
- Changed all data types to be a float

```
series = songs_df['Top Genre'].value_counts()
genreslessthan60 = series[series < 60]
songs_df = songs_df[~songs_df['Top Genre'].isin(genreslessthan60.index)]
```

```
1 print(songs_df.shape)
2 songs_df.info()
✓ 0.1s
```

```
(940, 11)
<class 'pandas.core.frame.DataFrame'>
Int64Index: 940 entries, 0 to 1993
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Year                  933 non-null   float64
 1   Beats Per Minute (BPM) 931 non-null   float64
 2   Energy                932 non-null   float64
 3   Danceability           464 non-null   float64
 4   Loudness (dB)          934 non-null   float64
 5   Liveness               933 non-null   float64
 6   Valence                433 non-null   float64
 7   Length (Duration)      932 non-null   float64
 8   Acousticness           933 non-null   float64
 9   Speechiness            937 non-null   float64
10   Top Genre              932 non-null   object
dtypes: float64(10), object(1)
```

	Year	Beats Per Minute (BPM)	Energy	Danceability	Loudness (dB)	Liveness	Valence	Length (Duration)	Acousticness	Speechiness	Top Genre
0	2004.0	157.0	30.0	53.0	-14.0	11.0	68.0	201.0	94.0	3.0	adult standards
1	2000.0	135.0	79.0	50.0	-11.0	17.0	81.0	207.0	17.0	7.0	album rock
3	2007.0	173.0	96.0	43.0	-4.0	3.0	37.0	269.0	0.0	4.0	alternative metal
10	2002.0	109.0	5.0	44.0	-16.0	11.0	31.0	162.0	88.0	4.0	adult standards
11	2003.0	124.0	46.0	74.0	-8.0	26.0	32.0	232.0	1.0	8.0	alternative rock

First Implementation (No Tuning)

```
# Model without tuning
dataset = songs_df.values

X = dataset[:,0:len(songs_df.columns)-1]
Y = dataset[:,len(songs_df.columns)-1]

X[X == '?'] = np.nan
X = X.astype(float)

label_encoded_Y = LabelEncoder().fit_transform(Y)

seed = 13
test_size = 0.25
X_train, X_test, Y_train, Y_test = train_test_split(X, label_encoded_Y, test_size = test_size, random_state = seed)

model = XGBClassifier()
model.fit(X_train, Y_train)
print(model)

predictions = model.predict(X_test)

accuracy = accuracy_score(Y_test, predictions)

print(f'Accuracy: {accuracy * 100.0:.2f}%')
```

Accuracy: 53.19%

Final Implementation (w/ Tuning)

Hyperparameters:

- max_depth
- learning_rate
- n_estimators

Accuracy: 63.83%

```
# Model with tuning
dataset = songs_df.values

X = dataset[:,0:len(songs_df.columns)-1]
Y = dataset[:,len(songs_df.columns)-1]

label_encoded_Y = LabelEncoder().fit_transform(Y)

seed = 7
test_size = 0.20
X_train, X_test, Y_train, Y_test = train_test_split(X, label_encoded_Y, test_size = test_size, random_state = seed)

model = XGBClassifier(max_depth=1, learning_rate=0.25, n_estimators=175)
model.fit(X_train, Y_train)
print(model)

predictions = model.predict(X_test)

accuracy = accuracy_score(Y_test, predictions)

print(f'Accuracy: {accuracy * 100.0:.2f}%')
```

Comparison (vs. Naive Bayes)

```
gnb = GaussianNB()

y_pred = gnb.fit(X_train, Y_train).predict(X_test)

num_mislabeled = ((Y_test != y_pred).sum() / X_test.shape[0]) * 100
print(f"Number of mislabeled points: {num_mislabeled:.2f}%")
print(f"Accuracy score: {100-num_mislabeled:.2f}%")
```

```
Number of mislabeled points: 54.79%
Accuracy score: 45.21%
```

Suggestions for Improvement

- More **data**
- More **representative selection** of the music genres
- Potentially **standardize** the data
- Using **gridsearch**